

NoSQL Databases with MongoDB

Surendra Panpaliya

2KNO₃ + H₂CO₃ → KCO₃ +

Founder and CEO, GKTC Innovations

<https://www.linkedin.com/in/surendrarp>



Surendra Panpaliya, DTM

**Founder & CEO, GKTCS Innovations –
building future-ready enterprises.**

**Empowered 35,000+ IT professionals
through training, mentoring, and
consulting.**

**Partnered with 300+ multinational
corporations to accelerate business
growth.**



Learning Objectives



Understand NoSQL fundamentals and



how they differ from RDBMS.



Model data efficiently for NoSQL systems.



Perform CRUD operations, aggregations, and



indexing in MongoDB.



Learning Objectives



Apply NoSQL patterns to real-world banking scenarios



e.g., Customer360, Transaction Logging, Fraud Detection



Build and query NoSQL databases



using MongoDB tools.



Prerequisites



Basic understanding of **Databases & SQL** concepts



tables, joins, normalization



Familiarity with **Python / JavaScript** syntax



for data access (optional but helpful).



Prior exposure to **JSON** and **REST APIs** recommended.



Lab Setup Requirements

System Requirements:

Laptop Configuration:

8 GB RAM minimum, 20 GB free disk space

Operating System:

Windows 10+, macOS, or Ubuntu



Lab Setup Requirements

Software Installation:

MongoDB Community Edition (latest version)

MongoDB Compass – GUI for query visualization

VS Code or **Jupyter Notebook** for scripting

Python 3.11+ with `pymongo` library installed

Agenda

Day1

Introduction to NoSQL & MongoDB Fundamentals

Day2

Data Modeling, Aggregation, and Use Cases

Module 1: NoSQL Overview



Evolution of Databases:



RDBMS → NoSQL



Key Characteristics:



Scalability, Flexibility,



Schema-less Design


Module 1: NoSQL Overview



**Types of NoSQL
Databases:**




Key-Value Stores
(Redis)



Document Stores
(MongoDB)



Column Stores
(Cassandra)



Graph Databases
(Neo4j)

Module 1: NoSQL Overview

Relational vs. NoSQL

Comparison table

Architecture diagram

Module 2: MongoDB Essentials



MongoDB Architecture:



Database → Collection → Document



BSON & JSON structure



CRUD Operations:



`insertOne()`, `find()`,



`updateOne()`, `deleteOne()`

Module 2: MongoDB Essentials

Indexing &
Query
Optimization

Data
Import/Export
using

`mongoimport`,
`mongoexport`

Module 3: Hands-On Lab



Lab 1: Create and Query a MongoDB Customer Database



Create a customers collection



Insert 10 customer profiles (JSON documents)



Query customers using filters (`find()`, `$and`, `$or`)



Update address & contact info using update operators

Day 2: Data Modeling, Aggregation, and Use Cases

Surendra Panpaliya

Global IT Trainer | Consultant | Thought Leader

Founder & CEO, GKTCS Innovations | DTM

Module 4: Data Modeling in NoSQL

Document Design Principles

Embedded vs. Referenced Documents

Denormalization vs. Normalization trade-offs

Module 4: Data Modeling in NoSQL



Schema Design for:



Customer Profiles



Account Transactions



Product Catalogs



Best Practices for
Scalability & Indexing

Module 5: Querying & Aggregations



Aggregation Pipeline Basics



`$match, $group, $project, $sort, $limit`



Filtering and transforming data



Case Study: Summarize total transactions
per customer

Module 6: Banking Use Cases & Labs



Lab 2: Implement Account Transaction Logging using MongoDB



Create a transactions collection



Insert transactions with timestamps, account numbers, and amounts



Query high-value transactions (`$gt`, `$sum`)



Build a fraud-detection query pipeline

Module 7: Capstone Assignment & Discussion



Design a Mini Customer360 Database for DBS Tech Bank



Collections: customers, accounts, transactions, alerts



Define key fields, indexing strategy, and relationships



Submit ER diagram (document-based) and sample queries

Module 4: Data Modeling in NoSQL

Document Design Principles

Embedded vs. Referenced Documents

Denormalization vs. Normalization trade-offs

What Is Data Modeling in NoSQL?



Designing how your data is stored, connected, and retrieved using:



Documents (instead of rows)



Collections (instead of tables)



Embedded data (nested JSON)



References (link documents)

What Is Data Modeling in NoSQL?

NoSQL is **schema-flexible**,

So design focuses on:

Performance

Query patterns

Real-world objects (Customer, Account, Transactions)

Core Document Design Principles (MongoDB)

Principle 1: Model Data Around Queries (Not Just Structure)

Design
documents

based on

how your app
queries data.

Example Query in DBS Bank

Fetch
Customer360

customer + all
accounts + balance
+ last 5
transactions

Embed account summary inside customer document:

```
{  
  "custId": "C1001",  
  "name": "Rohit Sharma",  
  "accounts": [  
    { "accNo": "SAV1001", "type": "Savings", "balance": 150000 },  
    { "accNo": "CRD2001", "type": "Credit Card", "limit": 100000 }  
  ]  
}
```

Why?



Only **1 query** to get full customer view.



Faster response for dashboards.

Principle 2: Embed When Data Has 1–N Relationship & Low Growth

Use embedded documents when child data is:

Small

Frequently read with parent


Not updated independently

Example: Customer → Address


A customer may have 1–2 addresses.

```
{  
  "custId": "C1001",  
  "name": "Rohit Sharma",  
  "address": {  
    "line1": "Pune Nagar Road",  
    "city": "Pune",  
    "pin": "411014"  
  }  
}
```


Why?



SIMPLE



FAST



ALL CUSTOMER INFO IN
ONE DOCUMENT

✗ When NOT to Embed

If child grows
indefinitely


If many
updates

If too large
(MongoDB
16MB limit)


✖ When NOT to Embed



Example: **transactions**



A customer may have
millions of transactions



DO NOT embed.

Principle 3: Reference When Data Grows Large (1–Many, Many–Many)

Example:

Customer → Transactions

Store transactions in a separate collection:

customers collection

```
{  
  "custId": "C1001",  
  "name": "Rohit Sharma"  
}
```

transactions collection

```
{  
  "txnId": "T9001",  
  "custId": "C1001",  
  "type": "DEBIT",  
  "amount": 2000,  
  "timestamp": "2025-01-01T10:00:00Z"  
}
```

How to fetch?


```
db.transactions.find({ custId: "C1001" })
```

Efficient


No large document issue

Transactions scalable (billions)

Principle 4: Use Denormalization for Speed



MongoDB prefers



duplicate data



to speed up reads.

Example

Store customer name in the transaction document.

```
{  
  "txnId": "T9001",  
  "custId": "C1001",  
  "customerName": "Rohit Sharma",  
  "amount": 2000  
}
```

Why?

Faster reporting

No join needed

Slight duplication
acceptable in
NoSQL

Principle 5: Use Aggregation Pipelines Instead of Joins

MongoDB aggregates are powerful.

Example: Customer + last transaction

```
db.transactions.aggregate([
  { $match: { custId: "C1001" } },
  { $sort: { timestamp: -1 } },
  { $limit: 1 }
])
```

Principle 6: Avoid Deep Nesting (> 3 Levels)

Bad (too deep):

```
{  
  "customer": {  
    "accounts": [  
      {  
        "transactions": [  
          { "notes": { "audit": { "updatedBy": "admin" }}}]  
        ]  
      }  
    ]  
  }  
}
```

Principle 6: Avoid Deep Nesting (> 3 Levels)



Good:



Keep customer, accounts, transactions separate



Link using custId, accNo

Principle 7: Use a Consistent Naming Pattern

Prefer

custId, accNo, txnId

lowercase + camelCase

Principle 7: Use a Consistent Naming Pattern



Avoid:



CUSTID, Cust_ID, CUST-id



This keeps documents clean and consistent.

Recommended MongoDB Model for DBS

customers

```
{  
  "custId": "C1001",  
  "name": "Rohit Sharma",  
  "mobile": "9876543210",  
  "addresses": [  
    { "type": "home", "city": "Pune", "pin": "411014" }  
  ]  
}
```

accounts

```
{  
  "accNo": "SAV1001",  
  "custId": "C1001",  
  "type": "Savings",  
  "balance": 150000  
}
```

transactions

```
{  
  "txnId": "T9001",  
  "accNo": "SAV1001",  
  "custId": "C1001",  
  "type": "DEBIT",  
  "amount": 2000,  
  "timestamp": "2025-01-01T10:00:00Z"  
}
```

How These Principles Help DBS Bank?

Requirement	Solution
Customer360	Embed account summary → fast queries
Millions of transactions	Reference in separate collection
Reporting dashboards	Denormalized fields
Audit & compliance	Aggregation pipelines
Scalability	No joins, simple relationships

1) Document Design Principles

Work backwards from access patterns

List top **reads, writes, filters, sorts, and aggregations.**

Co-locate fields that are **read together** → same document.

Use **compound indexes** that match your query prefixes.

Model by cardinality & growth

One-to-few:

embed (e.g., addresses in customer).

One-to-many (**bounded**):

embed or reference with subset embedding.

Model by cardinality & growth




**One-to-many
(unbounded/high-churn):**



Watch **unbounded arrays**



reference or **bucket**.



(can cause bloated docs & update contention).

Operational constraints




16 MB **document size** limit.



Prefer **immutable event docs** (append-only transactions).



Use **Decimal128** for money.



Add **createdAt / updatedAt / version** for audit & forward-compat.

Embedded vs. Referenced Documents



MONGODB GIVES TWO
WAYS TO CONNECT DATA:



**A. EMBEDDED
DOCUMENTS**



**B. REFERENCED
DOCUMENTS**

A. Embedded Documents

Store data INSIDE parent document

Use when:

Relationship is 1-to-few

Data is always needed together

Data size is small

Data rarely changes independently

Example (DBS Bank)

Customer → Address

```
{  
  "custId": "C1001",  
  "name": "Rohit Sharma",  
  "addresses": [  
    {  
      "type": "home",  
      "city": "Pune",  
      "pin": "411014"  
    },
```

Example (DBS Bank)

```
"pin": "411014"  
},  
{  
  "type": "work",  
  "city": "Mumbai",  
  "pin": "400093"  
}  
]  
}
```

Why embed?

Fetching customer profile

requires **only one query**

Faster reads for

Customer360 screen

When NOT to Embed

Do not embed when:

Child data grows without limit

Child data is frequently updated

Many users update same sub-document

Document may cross **16MB limit**

B. Referenced Documents

Use a separate collection + reference using ID

Use when:

1-to-many with large children

Data is updated independently

B. Referenced Documents

Use when:

You need scalability

Child collection is huge (millions)

Example (DBS Bank)

Customer → Transactions

```
customers
{
  "custId": "C1001",
  "name": "Rohit Sharma"
}
```

Example (DBS Bank)

Customer → Transactions

transactions

```
{  
  "txnid": "T9001",  
  "custId": "C1001",  
  "amount": 2000,  
  "type": "DEBIT"  
}
```

Example (DBS Bank)

Customer → Transactions



To fetch transactions:



```
db.transactions.find({ custId: "C1001" })
```

Why reference?

Transactions grow very fast

Updating one transaction should not

rewrite entire customer document

Better for analytics and compliance

Quick Comparison Table

Feature	Embedded	Referenced
Relationship	1-to-few	1-to-many, many-to-many
Read performance	 Fast	 Medium
Write performance	Medium	Fast for large data
Consistency	Strong	Medium
Document size risk	Yes	No
Recommended for	Customer profile, addresses	Accounts, transactions

Embedded vs. Referenced

Factor	Embed (same doc)	Reference (separate doc + id)	Banking example
Read locality	✓ Excellent	✗ Join at app/2nd query	Customer + top 3 addresses
Write frequency	✗ Hot if subdocs change often	✓ Isolate churn	KYC updates separate from profile
Cardinality	✓ One-to-few	✓ One-to-many/unbounded	Customer → many transactions

Embedded vs. Referenced

Factor	Embed (same doc)	Reference (separate doc + id)	Banking example
Reuse across parents	✗ Hard	✓ Easy	Shared branch data across customers
Size growth	✗ Risk of 16MB	✓ Bounded parent size	Large statements per month
Consistency need	✓ Single-doc atomicity	✗ Multi-doc consistency to manage	Profile + prefs together

2. Denormalization vs. Normalization Trade-offs

MongoDB
encourages

denormalization

(duplicate data)
but carefully.

A. Normalization (Like RDBMS)

Characteristics:

No duplicate data

Relations through references

Write operations easy

Read operations may require multiple queries


Mongo-style normalized data:

customers


```
{ "custId": "C1001", "name": "Rohit Sharma" }
```

- **transactions**
- { "txnId": "T9001", "custId": "C1001", "amount": 2000 }

When to use?



Very large relational
datasets



Frequently changing
fields



Transaction-heavy
systems

B. Denormalization (MongoDB Preferred)

Duplicate
some data

to avoid
joins and

speed
reads.

Example (DBS Bank)

Store **customerName** inside each transaction:

```
{  
  "txnId": "T9001",  
  "custId": "C1001",  
  "customerName": "Rohit Sharma",  
  "amount": 2000  
}
```

Why denormalize?

Faster dashboard queries

No JOIN required

Simpler and scalable

Good for analytics

✖️ Downsides of Denormalization

Problem	Example
Duplicate data	customerName stored in 10,000 transactions
Update overhead	If customer name changes → update all documents
More storage	Duplicate copies of same fields

Quick Trade-off Summary

Topic	Normalization	Denormalization
Read performance	✗ Slower	⭐ Fastest
Write performance	⭐ Fast	✗ Slower
Data consistency	⭐ Strong	✗ Weaker
Storage	⭐ Efficient	✗ Larger
Use-case	OLTP, updates	Analytics, dashboards

Best Practices for DBS Bank Domain

Use Embedding for

Addresses

KYC details

Preferences

Branch details (small dataset)

Best Practices for DBS Bank Domain



Use Referencing for:



Transactions



Accounts



Cards



Loans



Audit logs



Alerts/Notifications

Best Practices for DBS Bank Domain



Use Denormalization for:



Storing customerName, branchName, or category inside transactions



Faster Customer360



Faster reporting dashboards

Denormalize (duplicate small, stable fields)



Faster reads (fewer round trips); simpler queries.



Update fan-out (must update duplicates); risk of staleness.

Use when: read-heavy, fields change rarely (e.g.,
branchName).

Normalize (references)



✓ Single source of truth; smaller docs.



✗ Additional lookups/joins at app side.



Use when: write-heavy or shared/updating entities.

Safe denormalization pattern

Keep an **authoritative id**

branchId and **cache display fields**

(branchName) in the child.

Safe denormalization pattern

Periodically **refresh**

denormalized fields

via batch job or

change streams.


Schema Design Hands On



Customer Profiles




Account Transactions



Product Catalogs




Best Practices for
Scalability & Indexing



Module 5 – Querying & Aggregations

Aggregation Pipeline Basics



What is the Aggregation Pipeline?

- Like a **data processing assembly line**:
- Data flows through **stages**
- Each stage transforms the data
- Output of one stage → input to the next

Basic Syntax

```
db.collection.aggregate([
  { /* stage 1 */ },
  { /* stage 2 */ },
  { /* stage 3 */ }
])
```

For DBS Bank, suppose we have a transactions collection:

```
{  
  "txnId": "T9001",  
  "custId": "C1001",  
  "accNo": "SAV1001",  
  "type": "DEBIT",  
  "amount": 2000,  
  "channel": "UPI",  
  "txnDate": ISODate("2025-01-02T14:45:00Z")  
}
```

Find total DEBIT amount per customer
and show top 3 spenders.

\$match

\$group

\$project

\$sort

\$limit

Quick Cheat Sheet Table

Stage	Role	Typical Use in DBS Bank
\$match	Filter docs	Filter by year, type, channel
\$group	Aggregate/group	Total amount per custId/accNo
\$project	Shape fields, compute new fields	Rename, hide fields, compute averages
\$sort	Order results	Top spenders, latest txns
\$limit	Restrict number of rows	Top N customers / txns

Aggregation Pipeline Basics



Mental model



A pipeline is an **ordered list of stages**.



Each stage takes input documents and



outputs transformed documents



to the next stage.

Aggregation Pipeline Basics

Stage	Purpose (bank example)
\$match	Filter txns (e.g., last 30 days, DEBIT only)
\$project	Shape fields (rename, compute, hide)
\$group	Rollups (sum per account/customer)
\$sort	Order results (e.g., top spenders)

Aggregation Pipeline Basics

Stage	Purpose (bank example)
<code>\$limit</code>	Top-N
<code>\$addFields/\$set</code>	New fields (e.g., signed amount)
<code>\$lookup</code>	Join (e.g., txn → account → customer)
<code>\$unwind</code>	Flatten arrays (e.g., multiple addresses)
<code>\$facet</code>	Multi-result dashboards in one pass

Let's Connect



surendra@gktcs.com

[**https://www.linkedin.com/in/surendrarp**](https://www.linkedin.com/in/surendrarp)

[**https://www.gktcs.com**](https://www.gktcs.com)

Happy Learning@!!
Thanks for Your
Patience ☺

Surendra Panpaliya
GKTCS Innovations

